

Winter 2021 Data Science Intern Challenge

Question 1

a. To calculate the Average Order Value, the revenue must be divided by the number of orders. However, the 3,145.13 calculated is not a great reflection of this metric because it is merely the mean of the order amount.

A better evaluation of this data would be to find the average cost of the item being sold across all 100 stores. This will give Shopify a better sense of how much the product is priced.

b. The metric that will better represent this data is the average cost of the product.

```
In [69]: import pandas as pd
```

```
In [70]: df = pd.read_excel('/Data.xlsx')
```

```
In [71]: df.head()
```

```
Out[71]:
```

	order_id	shop_id	user_id	order_amount	total_items	payment_method	created_at
0	1	53	746	224	2	cash	2017-03-13 12:36:56.190
1	2	92	925	90	1	cash	2017-03-03 17:38:51.999
2	3	44	861	144	1	cash	2017-03-14 04:23:55.595
3	4	18	935	156	1	credit_card	2017-03-26 12:43:36.649
4	5	18	883	156	1	credit_card	2017-03-01 04:35:10.773

First, the data from the excel sheet is imported, and the information about this dataset is generated below. From there, we can tell that there is a total number of 5000 orders placed over a 30-day window.

In [72]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5000 entries, 0 to 4999
Data columns (total 7 columns):
#   Column          Non-Null Count  Dtype
---  -
0   order_id        5000 non-null   int64
1   shop_id         5000 non-null   int64
2   user_id         5000 non-null   int64
3   order_amount    5000 non-null   int64
4   total_items     5000 non-null   int64
5   payment_method  5000 non-null   object
6   created_at      5000 non-null   datetime64[ns]
dtypes: datetime64[ns](1), int64(5), object(1)
memory usage: 273.6+ KB
```

In [73]: `df['cost_per_item'] = df.order_amount/df.total_items`

To find the cost of each item sold in each shop, the order amount is divided by the total items for each order.

In [74]: `df.head()`

Out[74]:

	order_id	shop_id	user_id	order_amount	total_items	payment_method	created_at	cost_p
0	1	53	746	224	2	cash	2017-03-13 12:36:56.190	
1	2	92	925	90	1	cash	2017-03-03 17:38:51.999	
2	3	44	861	144	1	cash	2017-03-14 04:23:55.595	
3	4	18	935	156	1	credit_card	2017-03-26 12:43:36.649	
4	5	18	883	156	1	credit_card	2017-03-01 04:35:10.773	

From there, the average price for the product is calculated by dividing the sum of the cost by the total number of orders.

In [75]: `sum(df.cost_per_item)/len(df.order_id)`

Out[75]: 387.7428

c. 387.74 was the average cost found for the shoe, and considering the fact that these shops are relatively affordable, it seems to be on the high side.

```
In [84]: df1=df[['shop_id','cost_per_item']]  
df1.groupby(['shop_id']).mean()
```

Out[84]:


	cost_per_item
shop_id	
1	158.0
2	94.0
3	148.0
4	128.0
5	142.0
...	...
96	153.0
97	162.0
98	133.0
99	195.0
100	111.0

100 rows × 1 columns

It then became relevant to analyze the cost of the product across all 100 shops. There seemed to be a considerable disparity within the cost per item for each shop. This variation can be visualized in the box-plot below. The shop with ID 78 has a cost of 25725, making it the obvious outlier.

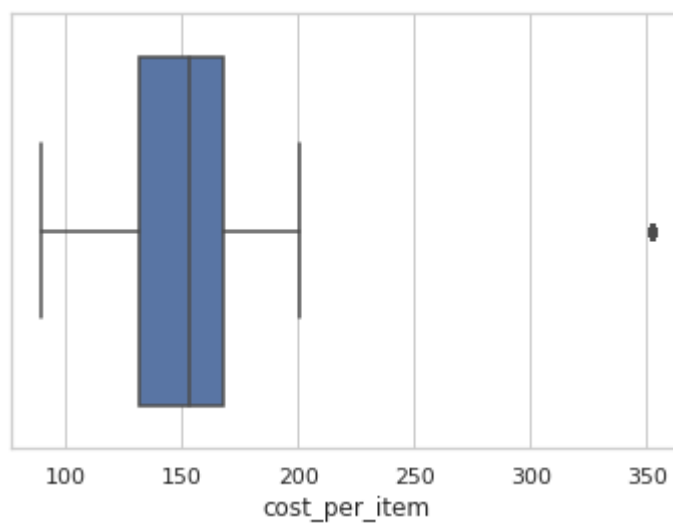
Indented block

```
In [85]: import seaborn as sns
sns.set(style="whitegrid")
ax = sns.boxplot(x=df1["cost_per_item"])
```



```
In [86]: df1 = df1[df1.shop_id != 78]
```

```
In [87]: ax = sns.boxplot(x=df1["cost_per_item"])
```



To better assess the data provided, the outlier was excluded. From the box plot above, the median is about 150, which is more reasonable for the item being sold.